

The background is a dark blue gradient with a series of concentric circles centered in the upper half. Overlaid on this are stylized, light blue circuit-like lines and nodes, primarily concentrated on the left and right edges, resembling a network or data flow diagram.




(MIS)CONFIGURING PAGE TABLES

BY ARTEM SHISHKIN




DISCLAIMER

I am presenting the contents of this presentation in my personal capacity. The views expressed are solely my own and do not reflect the views of Intel Corporation or its affiliates.





WHOAMI

- Security researcher at Intel Corporation
 - Intel STORM team member
 - Windows kernel enthusiast
- 



The background is a solid dark blue. In the four corners, there are decorative white line art elements resembling circuit traces or neural network connections. These lines are of varying thickness and connect to small white circles, creating a symmetrical, abstract pattern.

VIRTUAL MEMORY

FROM SCRATCH

VIRTUAL MEMORY CONCEPT

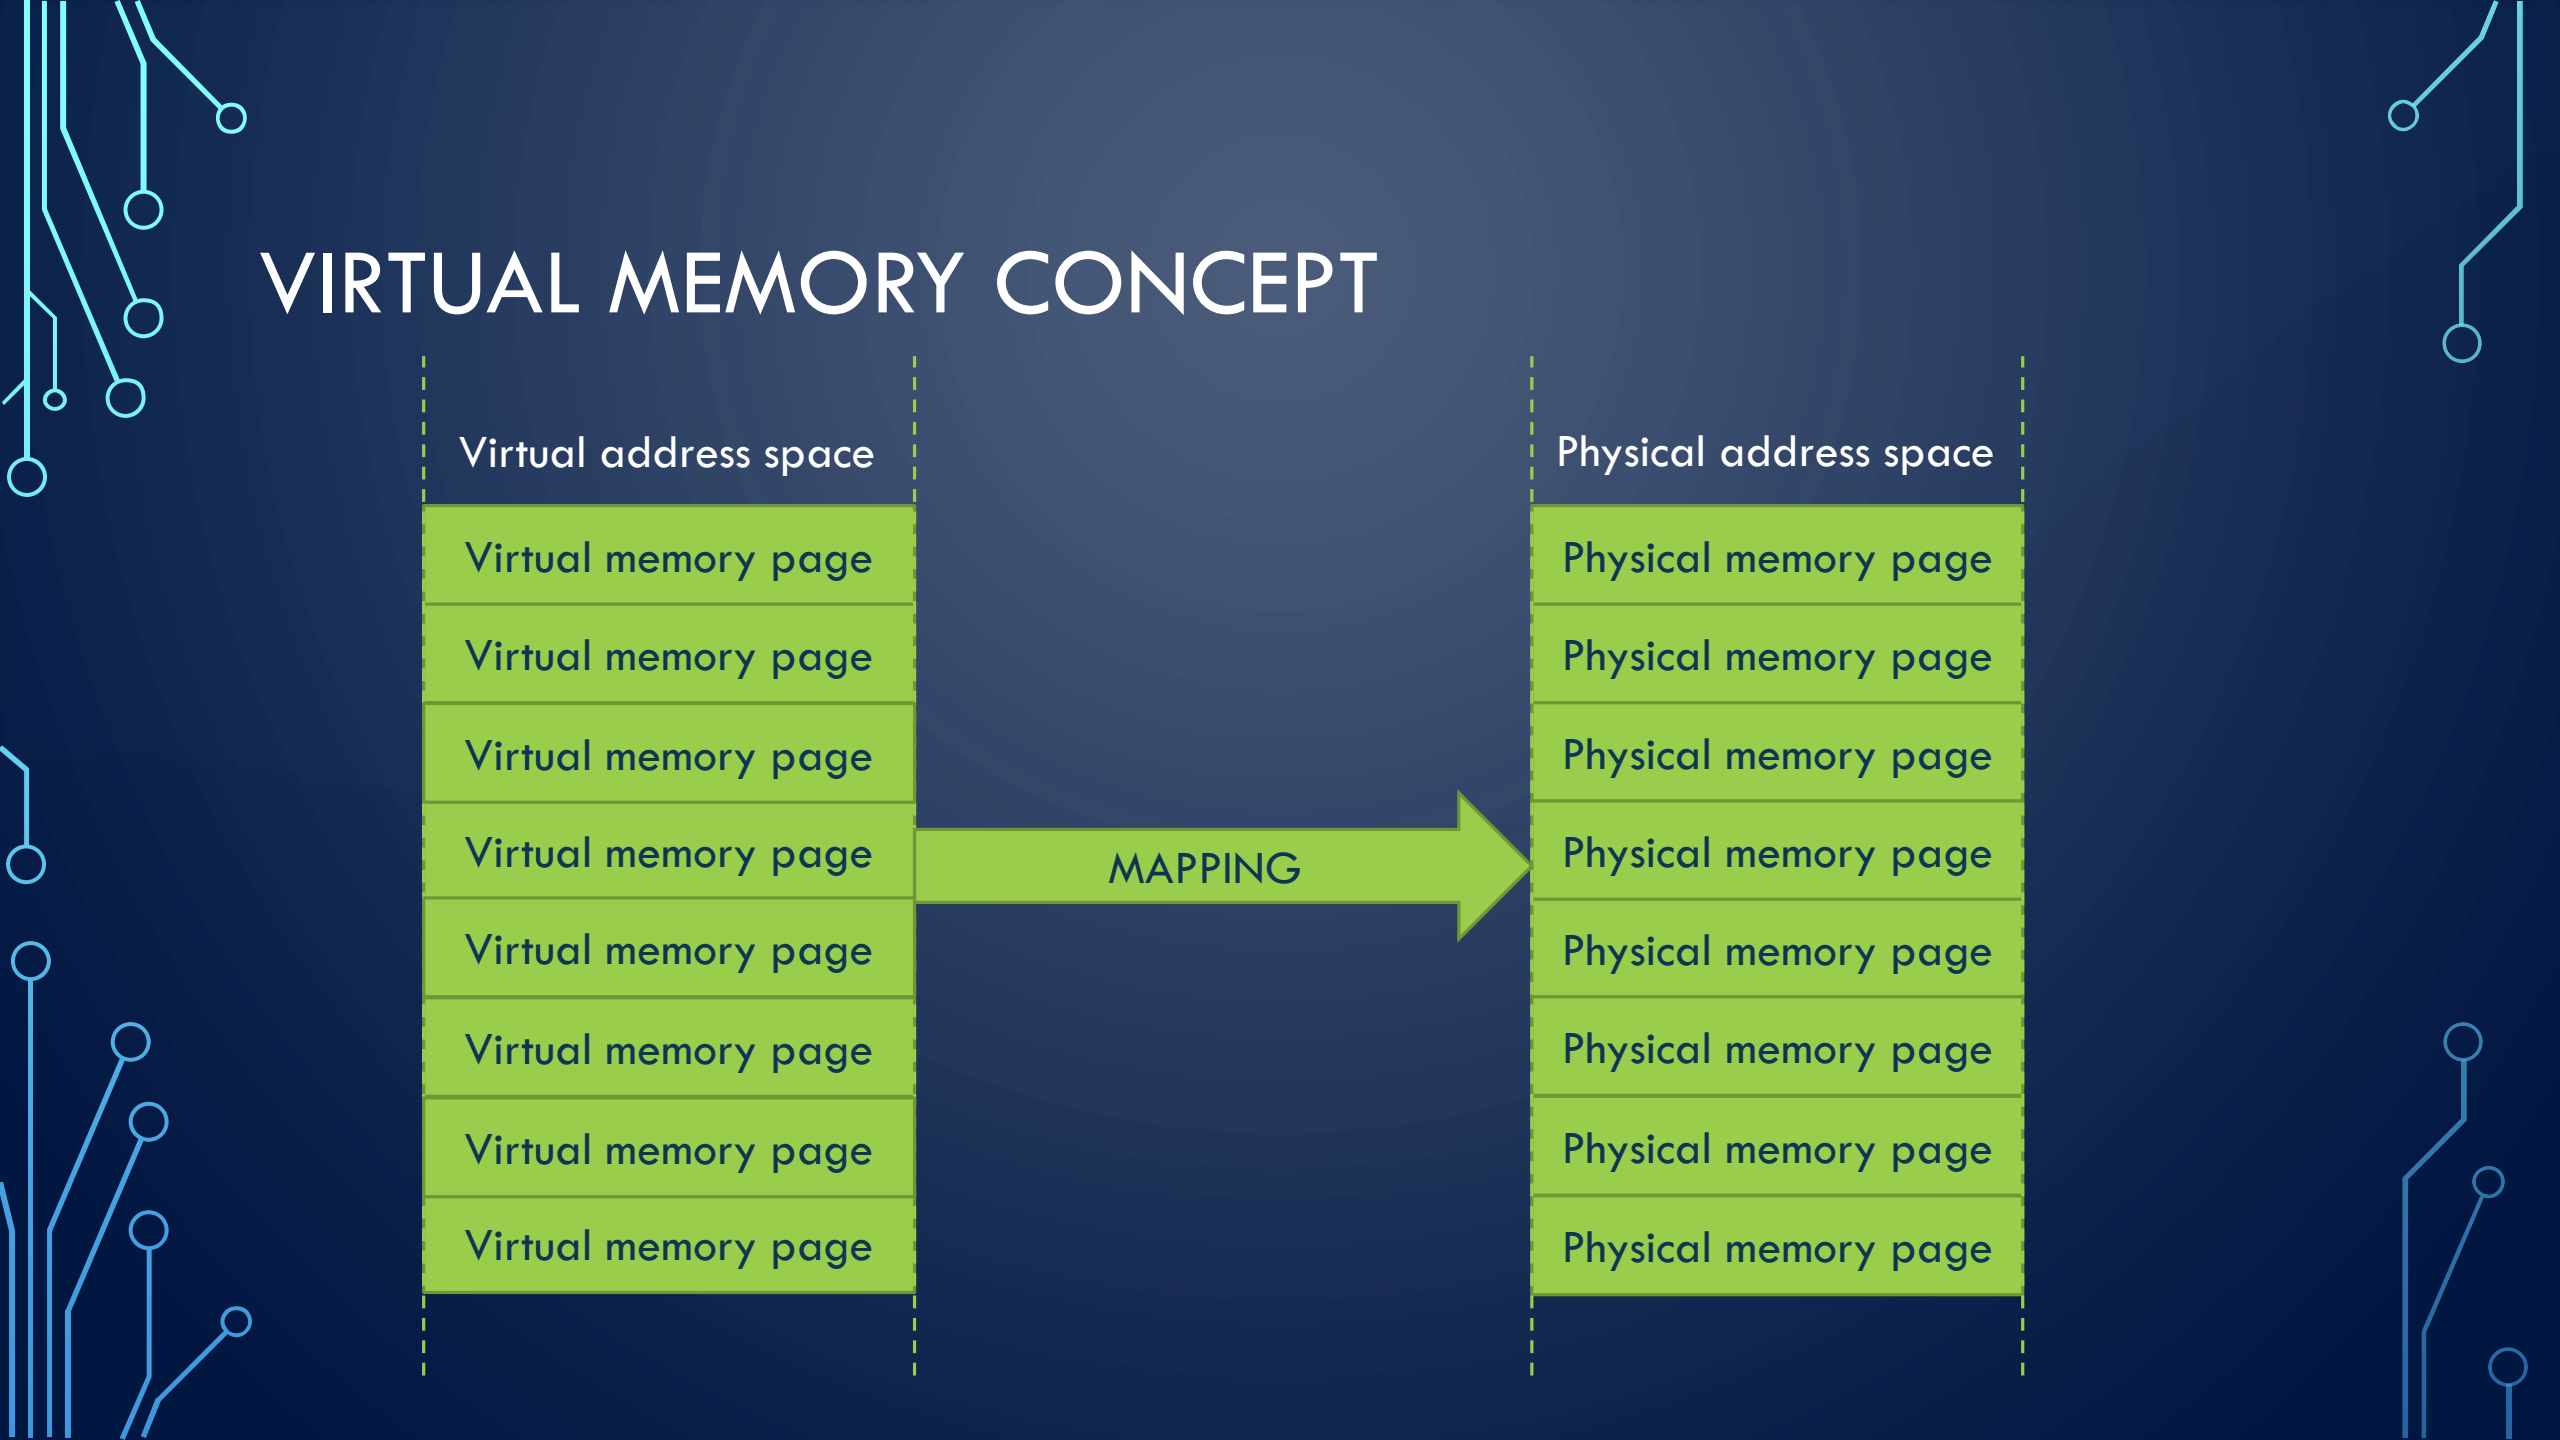
Virtual address space

Physical address space

MAPPING



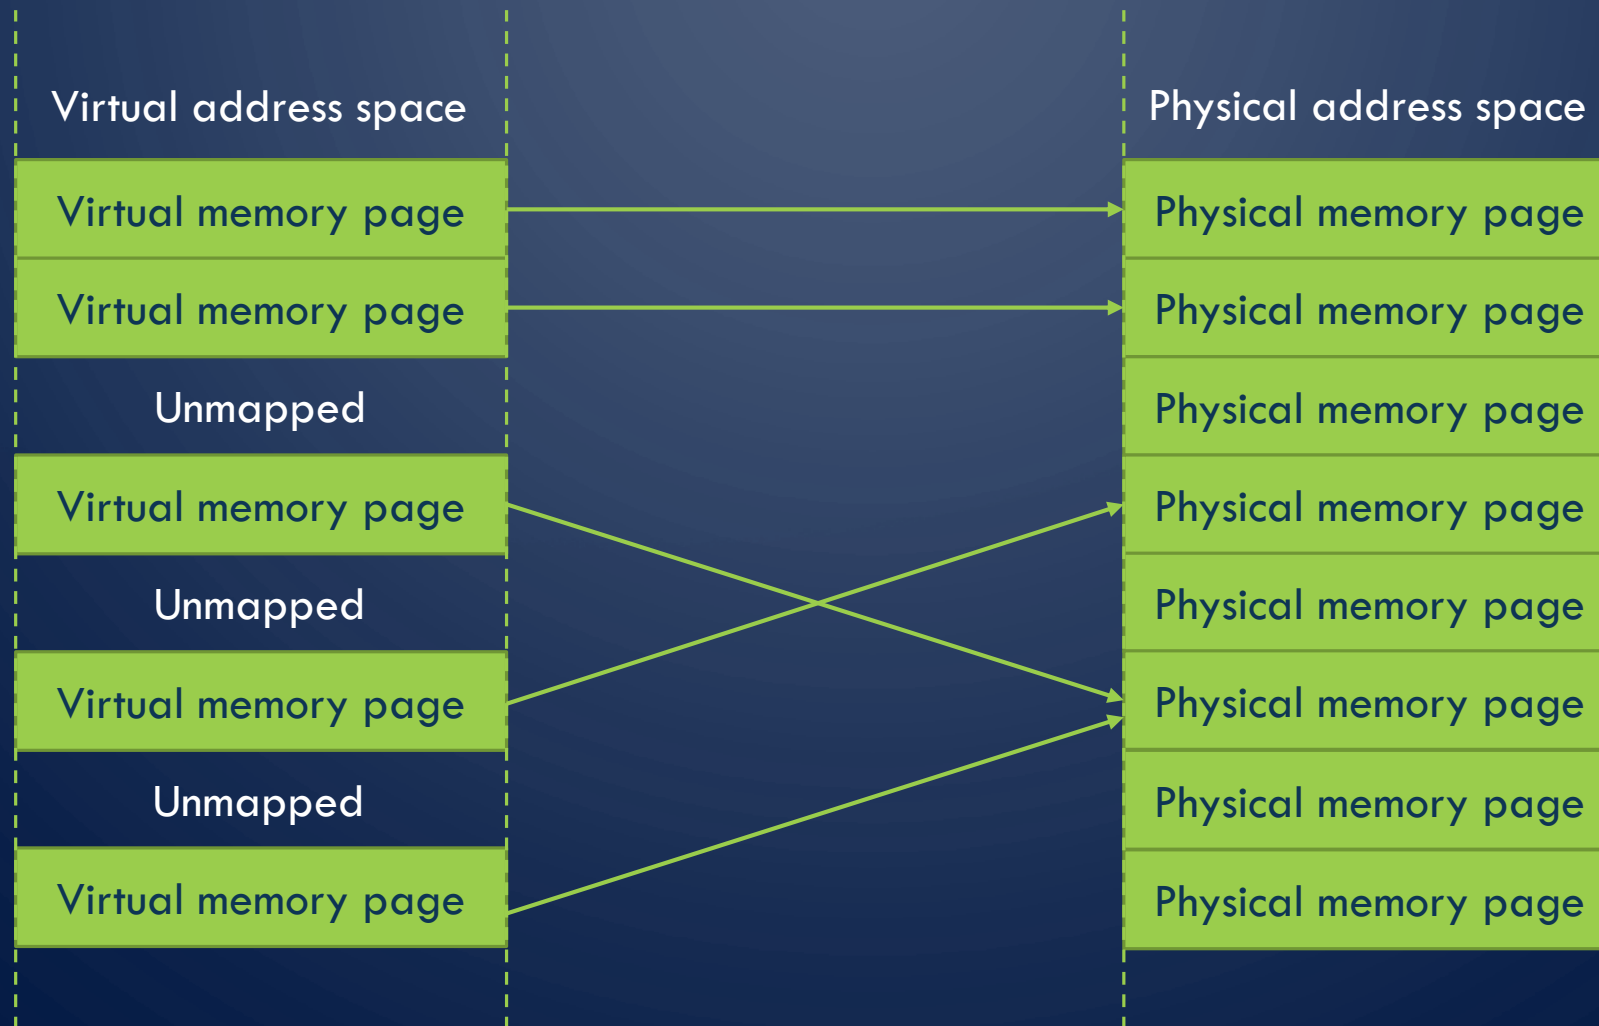
```
graph LR; V[Virtual address space] -- MAPPING --> P[Physical address space];
```

[illegible]

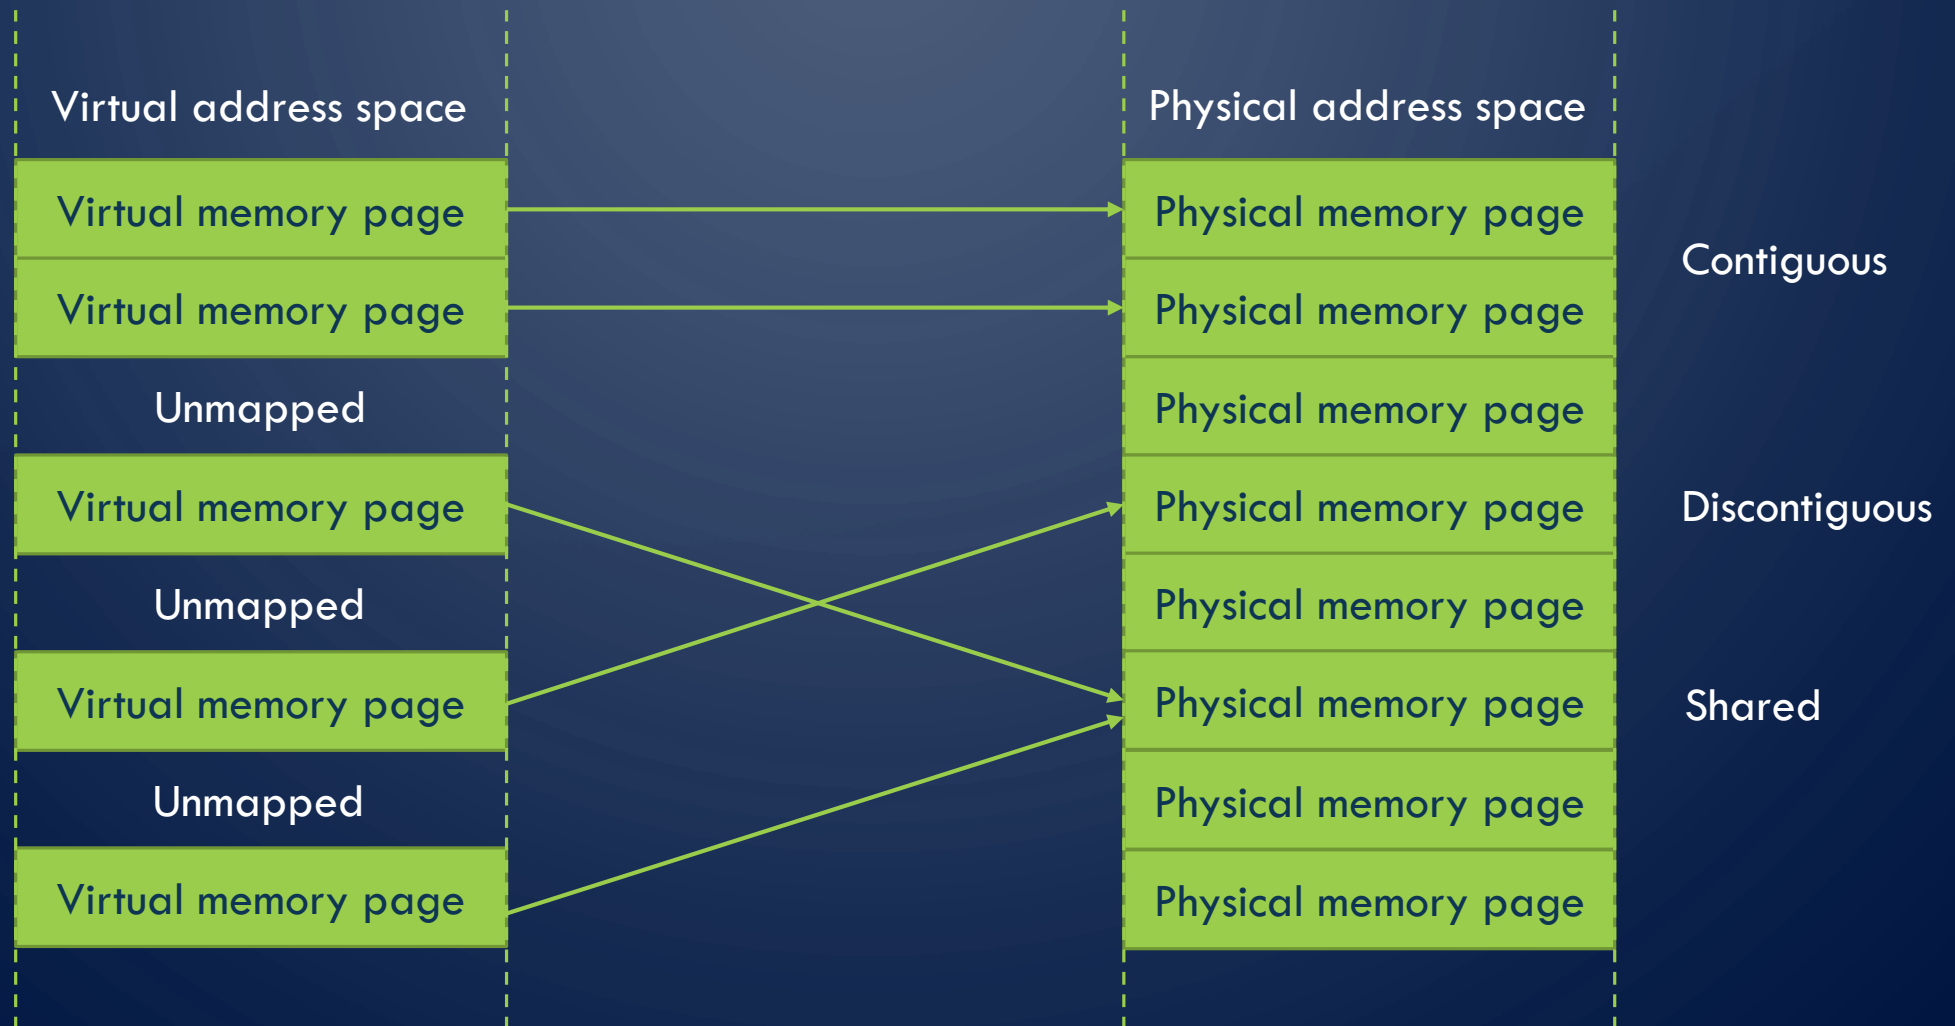
VIRTUAL MEMORY CONCEPT



VIRTUAL MEMORY CONCEPT



VIRTUAL MEMORY CONCEPT



VIRTUAL MEMORY BENEFITS

- Virtualize memory as a resource
 - Pretend we have a lot of it
 - Customize memory layout (like flat)
- Isolate execution environments
- Mitigate memory fragmentation
- Fine grained access control and protection
- A lot more, I'm just saying virtual memory is great

The background is a solid dark blue. In the four corners, there are abstract, light blue line art designs that resemble circuit traces or stylized trees. These designs consist of thin lines that branch out and terminate in small circles. The top-left and bottom-left designs are more complex and dense, while the top-right and bottom-right designs are simpler and more sparse.

IMPLEMENTATION

IT'S ALWAYS ABOUT THE DETAILS

X86-64, LONG MODE PAGING

- Implemented as 4 level page tables (48 bit virtual address)
 - 5 level available recently for long mode (56 bit virtual address)
- Root is stored in a CR3 register
- Pages can map several memory chunk sizes depending on CPU capabilities
 - 1 gigabytes
 - 2 megabytes
 - 4 kilobytes

X86-64, LONG MODE PAGING

CR3 0x12A000

VA 0xFFFFF80346490123

16 bit

9 bit

9 bit

9 bit

9 bit

12 bit

1111 1111 1111 1111 1111 1000 0000 0011 0100 0110 0100 1001 0000 0001 0010 0011

Canonical part (0s or 1s)


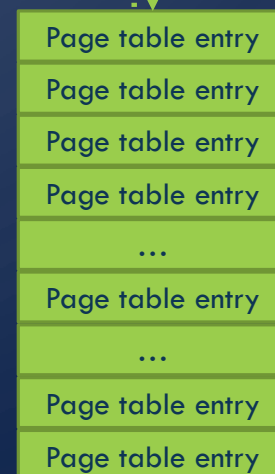
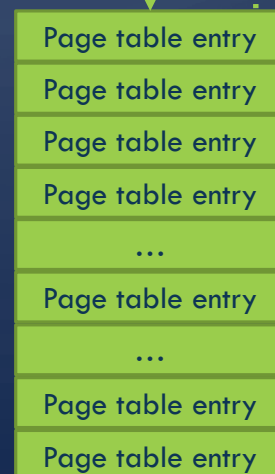
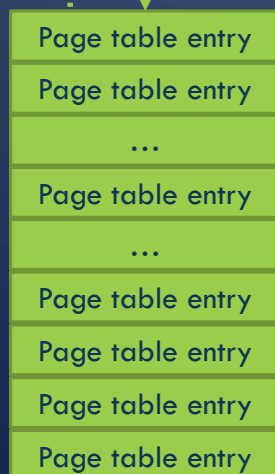
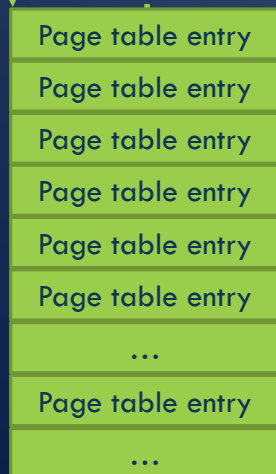
PML4 table
index
(0x1F0)

PDP table
index (0x0D)

Page
directory
index (0x32)

Page table
index (0x90)

Byte offset within the
page
(0x0123)



X86-64, LONG MODE PAGING

CR3 0x12A000

VA 0xFFFFF80346490123

16 bit

9 bit

9 bit

9 bit

21 bit

1111 1111 1111 1111 1111 1000 0000 0011 0100 0110 0100 1001 0000 0001 0010 0011


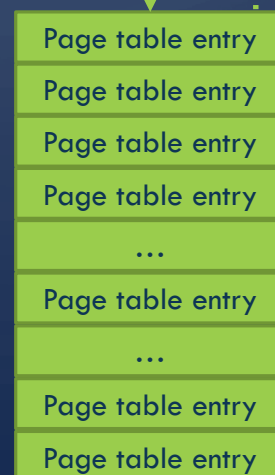
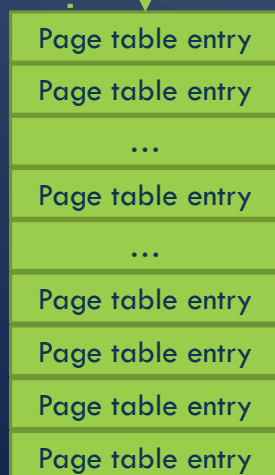
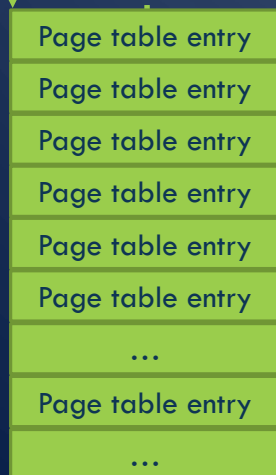
Canonical part (0s or 1s)

PML4 table
index
(0x1F0)

PDP table
index (0x0D)

Page
directory
index (0x32)

Byte offset within the page
(0x90123)



X86-64, LONG MODE PAGING

CR3 0x12A000

VA 0xFFFFF80346490123

16 bit

9 bit

9 bit

30 bit


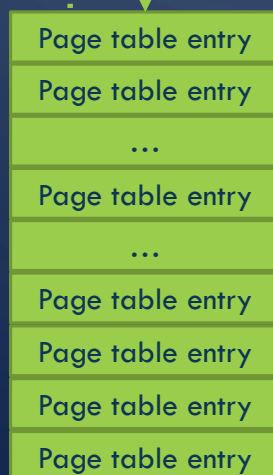

1111 1111 1111 1111 1111 1000 0000 0011 0100 0110 0100 1001 0000 0001 0010 0011

Canonical part (0s or 1s)

PML4 table
index
(0x1F0)

PDP table
index (0x0D)

Byte offset within the page
(0x46490123)



PAGE TABLE ENTRY

- Almost the same for every page table level
- Points to the next page table or a final physical page
- Stronger security attributes control the whole region described by the page table
 - Like R/W bit, disallowing writes for the whole 512 GB region if unset in PML4E
 - Or U/S bit, disallowing usermode access for the whole 1 GB region if unset in PDPTE
 - Or XD bit, disallowing instruction fetches from 2 MB region if set in PDE

PAGE TABLE ENTRY

X D	Prot. Key	Ignored	Rsvd.	Address of a page frame	Ign.	G	P A T	D	A	P C D	P W T	U / S	R / W	P	PTE: 4KB page
--------	--------------	---------	-------	-------------------------	------	---	-------------	---	---	-------------	-------------	-------------	-------------	---	---------------------

P – Present

R/W – Read / Write

U/S – User / Supervisor

PWT – Page Level Write Through

PCD – Page Level Cache Disable

A – Accessed

D – Dirty

PAT – Page Attribute Table

G – Global

XD – Execute disable

MAPPING THE PAGE

X D	Prot. Key	Ignored	Rsvd.	Address of a page frame	Ign.	G	P A T	D	A	P C D	P W T	U / S	R / W	P	PTE: 4KB page
--------	--------------	---------	-------	-------------------------	------	---	-------------	---	---	-------------	-------------	-------------	-------------	---	---------------------

P – Present

R/W – Read / Write

U/S – User / Supervisor

PWT – Page Level Write Through

PCD – Page Level Cache Disable

A – Accessed

D – Dirty

PAT – Page Attribute Table

G – Global

XD – Execute disable

PROTECTION

X D	Prot. Key	Ignored	Rsvd.	Address of a page frame	Ign.	G	P A T	D	A	P C D	P W T	U / S	R / W	P	PTE: 4KB page
--------	--------------	---------	-------	-------------------------	------	---	-------------	---	---	-------------	-------------	-------------	-------------	---	---------------------

P – Present

R/W – Read / Write

U/S – User / Supervisor

PWT – Page Level Write Through

PCD – Page Level Cache Disable

A – Accessed

D – Dirty

PAT – Page Attribute Table

G – Global

XD – Execute disable

CACHING

X D	Prot. Key	Ignored	Rsvd.	Address of a page frame	Ign.	G	P A T	D	A	P C D	P W T	U / S	R / W	P	PTE: 4KB page
--------	--------------	---------	-------	-------------------------	------	---	-------------	---	---	-------------	-------------	-------------	-------------	---	---------------------

P – Present

R/W – Read / Write

U/S – User / Supervisor

PWT – Page Level Write Through

PCD – Page Level Cache Disable

A – Accessed

D – Dirty

PAT – Page Attribute Table

G – Global

XD – Execute disable

IGNORED VS RESERVED

X D	Prot. Key	Ignored	Rsvd.	Address of a page frame	Ign.	G	P A T	D	A	P C D	P W T	U /S	R / W	1	PTE: 4KB page
Ignored														0	PTE: not present

- Reserved – not to be modified
- Ignored – feel free to modify
- Software PTEs for memory manager
- Working set hash table index

SCANNING PAGE TABLES

- I suggest relying on hardware PTEs
 - This ensures mapping working in hardware
 - Skip memory manager semantics
- Page tables walk
 - Like a tree walk
- Parse PTEs properly
 - Check the page size
 - Check if the page is mapped




STATIC PAGE TABLE ANALYSIS

WHAT DO WE LOOK FOR



USER VS KERNEL



USER VS KERNEL

- Check if U/S bit is set correctly
- Do you think it is a naive idea?
 - Well, check [CVE-2018-1038](#), a single incorrectly set bit led to overwriting page tables from usermode

W^X

- Check if the page is writable and executable at the same time
- Useful for the latest Windows kernel exploits
- Currently some memory is mapped as WX
 - UEFI runtime services
 - Part of ntoskrnl image
 - Custom driver allocations
 - Drivers that use legacy MmMapIoSpace (vs MmMapIoSpaceEx)
- Virtualization-based security prevents W^X at hypervisor level using extended page tables

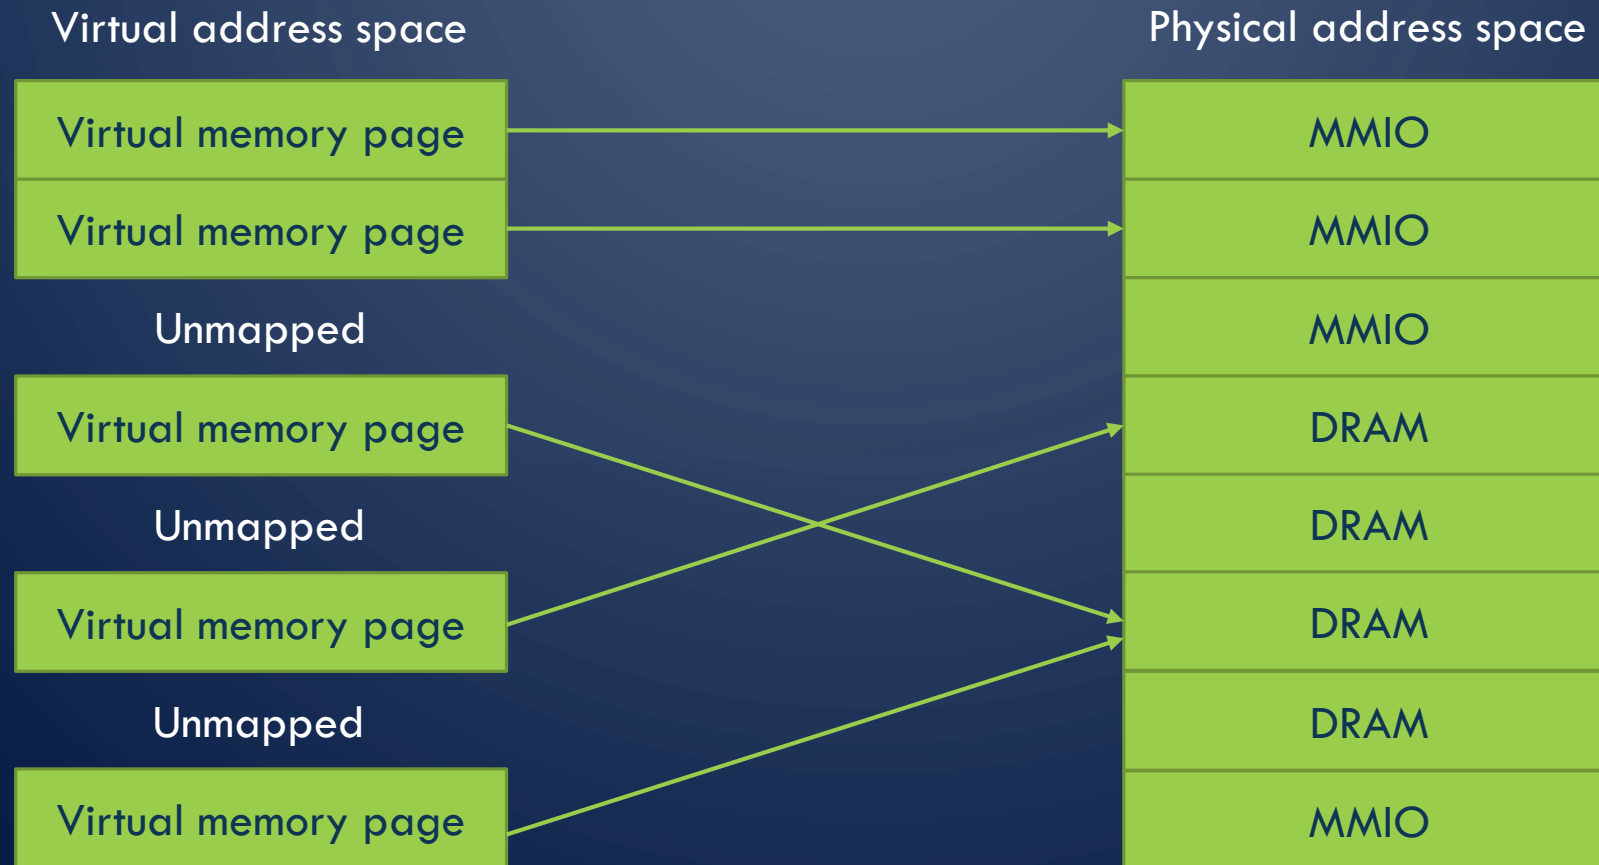
UEFI RANGE

- Basically mapped as RWX
 - Depends on a firmware
 - Modern MS Surface systems support page protection
- Function pointer table at `hal!HalEfiRuntimeServicesBlock`
 - Can be triggered by reading UEFI variable for example
- Blocked by VBS

SEMANTIC GAP

- `PVOID MmMapIoSpace(PHYSICAL_ADDRESS PhysicalAddress, SIZE_T NumberOfBytes, MEMORY_CACHING_TYPE CacheType);`
- `PVOID MmMapIoSpaceEx(PHYSICAL_ADDRESS PhysicalAddress, SIZE_T NumberOfBytes, ULONG Protect);`
- `MmMapIoSpace` goes to `MmMapIoSpaceEx`
 - Caching parameter implicitly converted to Protection parameter
 - `MmNonCached` converted to `PAGE_NO_CACHE | PAGE_EXECUTE_READWRITE`
 - `MmCached` converted to `PAGE_EXECUTE_READWRITE`
 - `MmWriteCombined` converted to `PAGE_WRITECOMBINE | PAGE_READWRITE`

ACCESS TO THE HARDWARE



ACCESS TO THE HARDWARE

- Check caching settings for the mapping
 - It is common that MMIO is mapped as non-cached
- Drivers can expose MMIOs to unprivileged code
- MMIO controls the hardware directly
- Immediate LPE if DMA engine is exposed
- <https://access.redhat.com/security/cve/cve-2010-5313>

SHARED MAPPINGS

- Cross-process
- Cross-mode
- Check if less privileged mode has higher privileged mapping
 - User can access sensitive kernel data, which leads to info leak
 - User can unexpectedly modify privileged mapping data, which leads to LPE
- Example: KUSER_SHARED_DATA
 - Read-only for user mode, Read-Write for kernel mode, no sensitive info



DYNAMIC PAGE TABLE ISSUES

WATCH THE CACHES



TLB AND PAGING STRUCTURE CACHES

- TLB - Translation lookaside buffers
- Per CPU thread
- Cache translation information
 - PFN address
 - R/W bit
 - U/S bit
 - XD bit
 - Protection keys
 - Dirty flag
 - Memory type

SHARED RESOURCE

- Page tables occupy a lot of memory
 - Mapping 16 GB of memory requires ~ 32 MB of memory to map as 4KB pages
 - Often shared between processes / modes
- We have a publicly documented OS APIs to change mapping attributes
- Memory manager is supposed to synchronize memory view across the CPU threads

SHARED RESOURCE

- Synchronization failure leads to race conditions
 - Example: CVE-2018-18281
- Check if caches are flushed and synchronized
 - Not trivial
 - But worth it

REFERENCES

- Intel Software Developers Manual
 - <https://software.intel.com/en-us/articles/intel-sdm>
- Can you break Windows Page Table Isolation?
 - <https://nadav.amit.zone/windows/2018/09/15/windows-pti.html>
- Taking a page from the kernel's book: A TLB issue in mremap()
 - <https://googleprojectzero.blogspot.com/2019/01/taking-page-from-kernels-book-tlb-issue.html>
- Total meltdown
 - <http://blog.frizk.net/2018/03/total-meltdown.html>

The background is a dark blue gradient. In the corners, there are decorative white line art elements resembling circuit boards or neural networks, with lines and small circles.

THANKS AND HAPPY FISHING!

QUESTIONS?